

# Making Reinforcement Learning Practical For Real World Developers

Praveen Veerath  
Sr. AI Solutions Architect  
Amazon Web Services

# THE AWS ML STACK

Broadest and deepest set of capabilities












## AI Services

VISION			SPEECH		LANGUAGE		CHATBOTS	FORECASTING	RECOMMENDATIONS
									
REKOGNITION IMAGE	REKOGNITION VIDEO	TEXTRACT	POLLY	TRANSCRIBE	TRANSLATE	COMPREHEND & COMPREHEND MEDICAL	LEX	FORECAST	PERSONALIZE


## ML Services

 Amazon SageMaker	Ground Truth	Notebooks	Algorithms + Marketplace	Reinforcement Learning	Training	Optimization	Deployment	Hosting
--	--------------	-----------	--------------------------	------------------------	----------	--------------	------------	---------

## ML Frameworks + Infrastructure

FRAMEWORKS	INTERFACES	INFRASTRUCTURE						
 TensorFlow	 mxnet		GLUON					
PYTORCH	 Keras							
		EC2 P3 & P3DN	EC2 G3	EC2 C5	FPGAS	GREENGRASS	ELASTIC INFERENCE	INFERTIA

# Reinforcement Learning in the broader AI context



# Machine Learning Overview




## **SUPERVISED**

Example Driven  
Training – every  
datum has a  
corresponding label



## **UNSUPERVISED**

No labels for  
training data




## **REINFORCEMENT**

Learns through  
consequences of  
actions in a specific  
environment

How do you build machine learning models that can make decisions when there is no labeled data?

# Complexity vs Data



# How do you learn in RL?

Learn by  
interacting with  
the  
environment

Trial and error

Observe  
results

Optimize  
learning  
strategy to  
maximize long  
term reward

Model learns  
how to make  
complex  
decisions

# What is an RL environment?

Real-world or a  
representation of  
the real world

Programmed to  
represent real  
world conditions

Enables  
interaction with  
user or a  
computer program

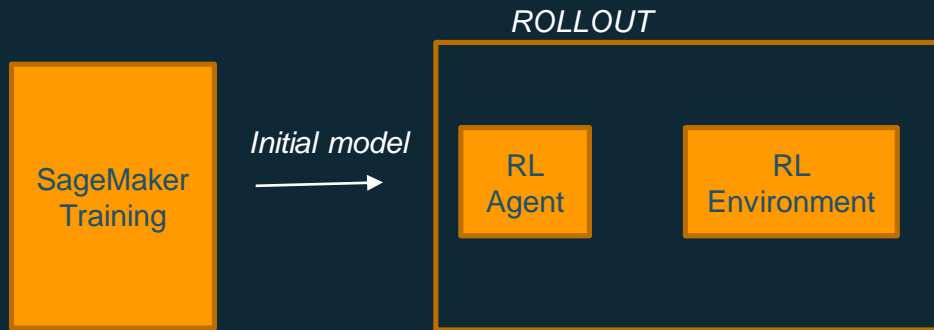
Dynamic and  
updates itself  
based on the  
interactions and  
programmed  
behavior



## How SageMaker RL works

*Let's train our humanoid, Harry, to walk*

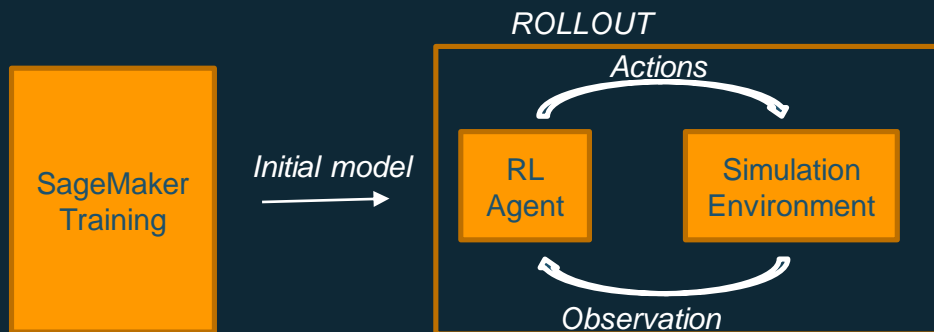
# The players




# At first, Harry can't even stand up



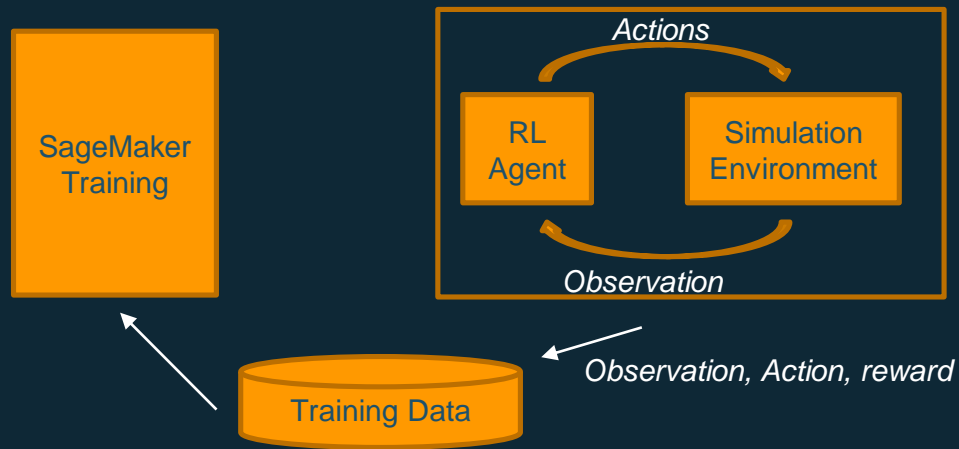
# Actions and Observations



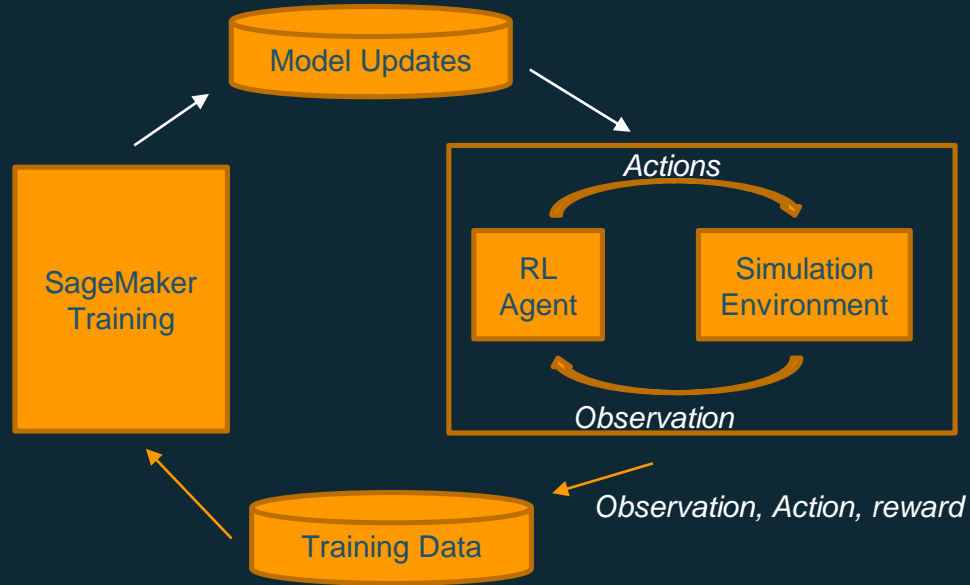
# Learning from Harry's actions and observations



# Interactions in the environment generate training data



# Training results in model updates

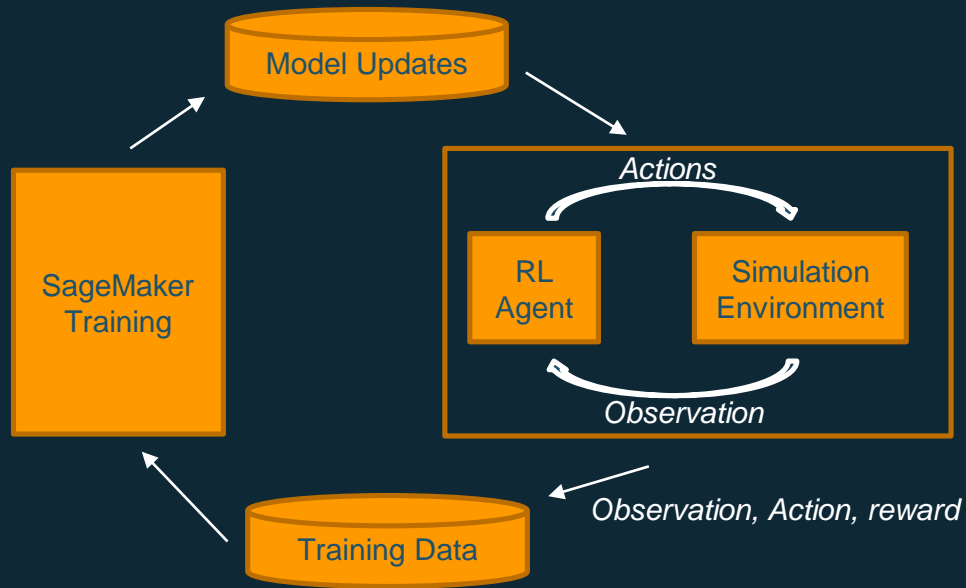


# Harry learns to stand and step







# Multiple training episodes improve learning



# RL Agents try to maximize rewards



# After many episodes of training

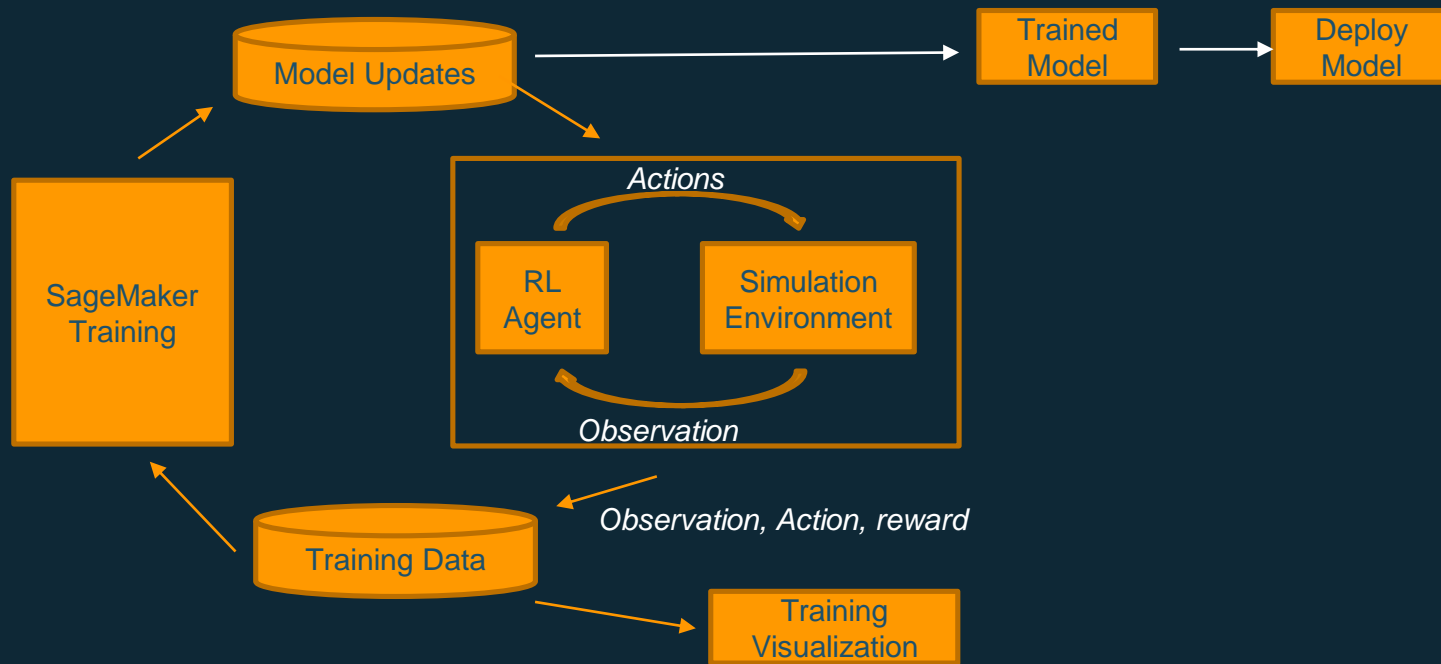


# Eventually, Harry learns how to walk and run



*You can continue training Harry to jump obstacles, play games, dance, and more*

# Evaluate and deploy trained models




# AWS DeepRacer Origin



How can we put  
Reinforcement  
Learning  
in the hands of all  
developers? *literally*

Our Objective: Teach developers RL

# AWS DeepRacer Car Specifications



**Car:** 1/18th scale 4WD with monster truck chassis

**CPU:** Intel Atom™ Processor

**Memory:** 4GB RAM

**Storage:** 32GB (expandable)

**Wi-Fi:** 802.11ac

**Camera:** 4 MP camera with MJPEG

**Software:** Ubuntu OS 16.04.3 LTS, Intel® OpenVINO™ toolkit, ROS Kinetic

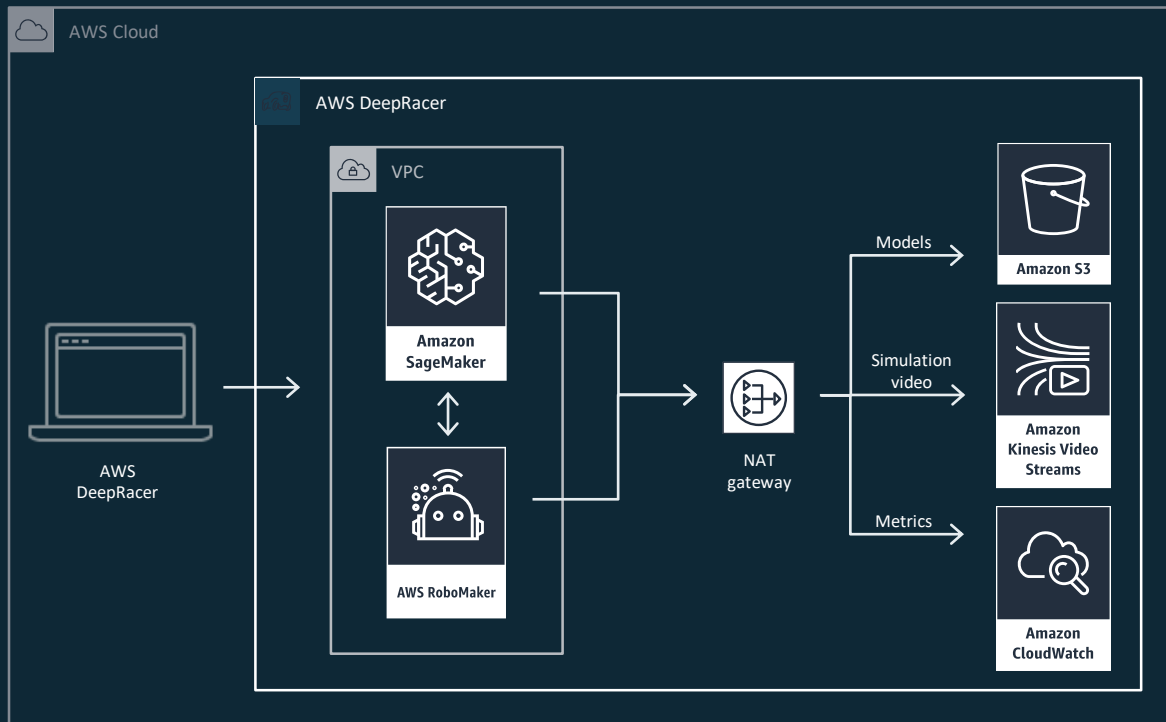
**Drive battery:** 7.4V/1100mAh lithium polymer

**Compute battery:** 13600mAh USB-C PD

**Ports:** 4x USB-A, 1x USB-C, 1x Micro-USB, 1x HDMI

**Sensors:** Integrated accelerometer and gyroscope

# AWS DeepRacer Architecture





# Reinforcement Learning with DeepRacer



**MODEL**



**AGENT**



**ACTION**




**ENVIRONMENT**



**GOAL**

# Coordinate system and track waypoints




# How to train a reinforcement learning model.

## Training an RL model

Training is an iterative process. In a simulator the agent explores the environment and builds up experience. The experiences collected are used to update the neural network periodically and the updated models are used to create more experiences.

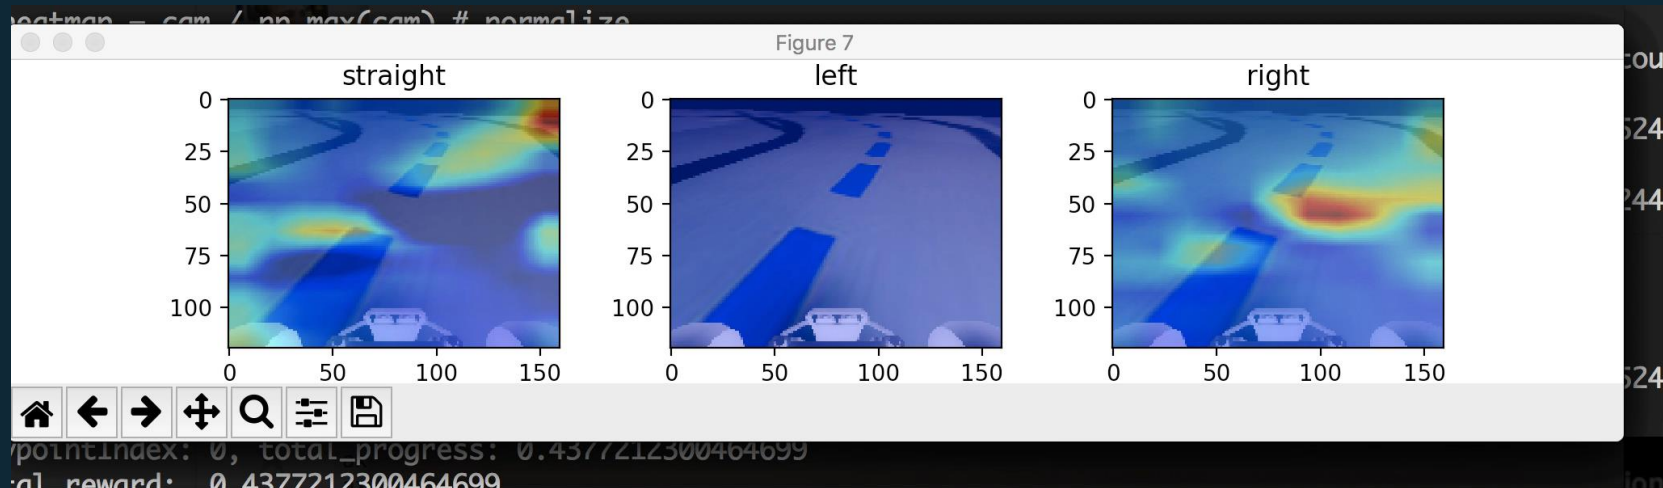
With AWS DeepRacer, we are training a vehicle to drive itself. It can be tricky to visualize the process of training, so let's take a look at a simplified example.

[Previous](#)[Next](#)

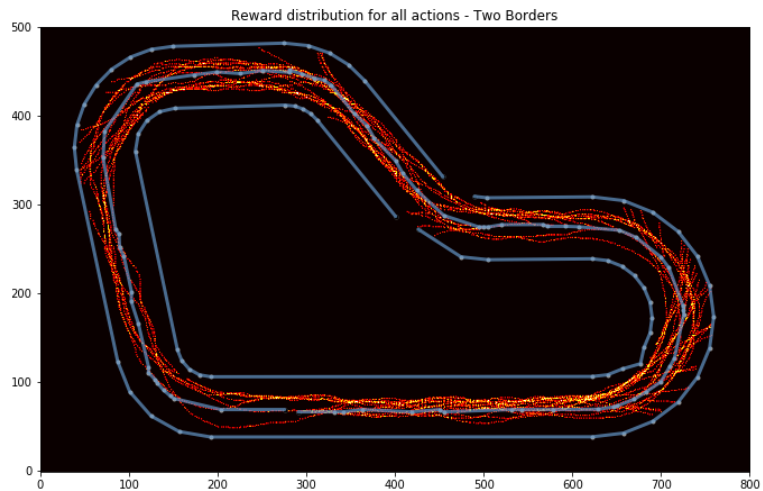
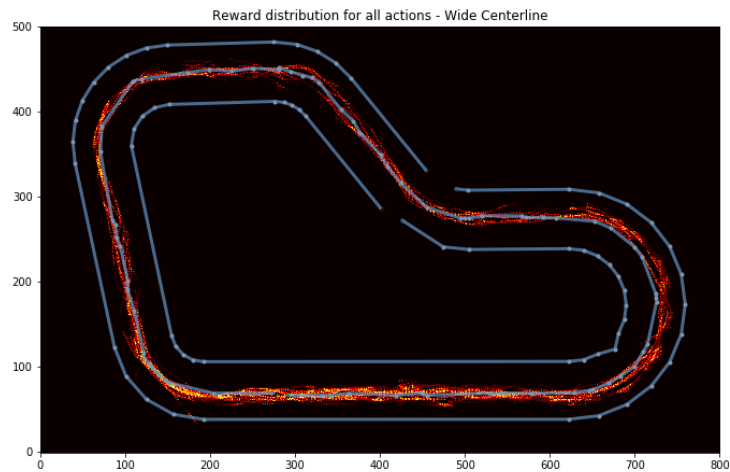
# Reward Function Example

```
def reward_function(params):  
    '''  
    Example of rewarding the agent to follow center line  
    '''  
  
    # Read input parameters  
    track_width = params['track_width']  
    distance_from_center = params['distance_from_center']  
  
    # Calculate 3 markers that are at varying distances away from the center line  
    marker_1 = 0.1 * track_width  
    marker_2 = 0.25 * track_width  
    marker_3 = 0.5 * track_width  
  
    # Give higher reward if the car is closer to center line and vice versa  
    if distance_from_center <= marker_1:  
        reward = 1.0  
    elif distance_from_center <= marker_2:  
        reward = 0.5  
    elif distance_from_center <= marker_3:  
        reward = 0.1  
    else:  
        reward = 1e-3 # likely crashed/ close to off track  
  
    return float(reward)
```




# Understand the model: Grad-CAM





# Reward debugging



# Reward Function – Center Lane




# Reward Function – Left Lane







# Sim2Real Debug




Optimal Path

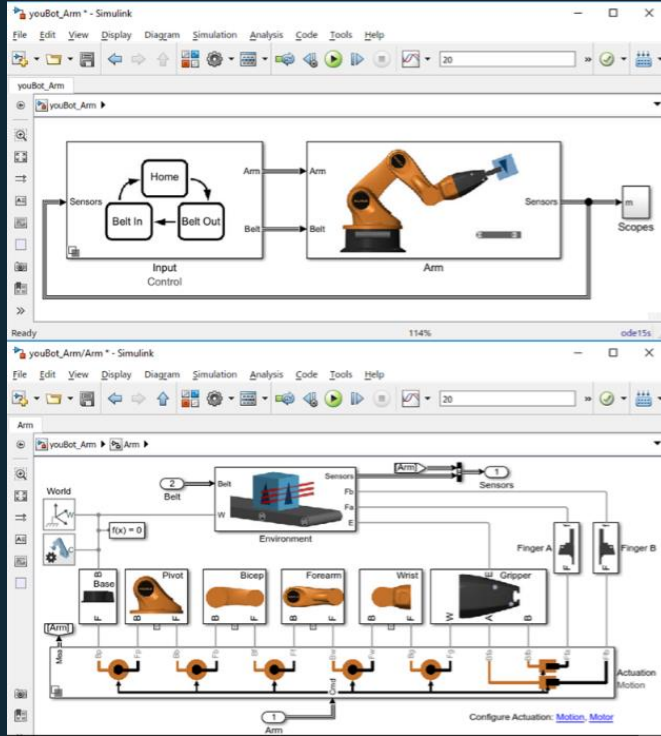


Converged path


# SageMaker RL + AWS RoboMaker accelerate learning for AWS DeepRacer




# You can use commercial simulation environments



MATLAB® and Simulink® for modeling and simulation





# Customers are using Amazon SageMaker RL



GE Healthcare

**HONDA**



mixi

amazon

SyntheticGestalt

Scientific Research by Artificially Intelligent Agents

 Tradelegs

# DeepRacer Links

DeepRacer Simulator:

<https://github.com/aws-robotics/aws-robomaker-sample-application-deepracer>

SageMaker RL:

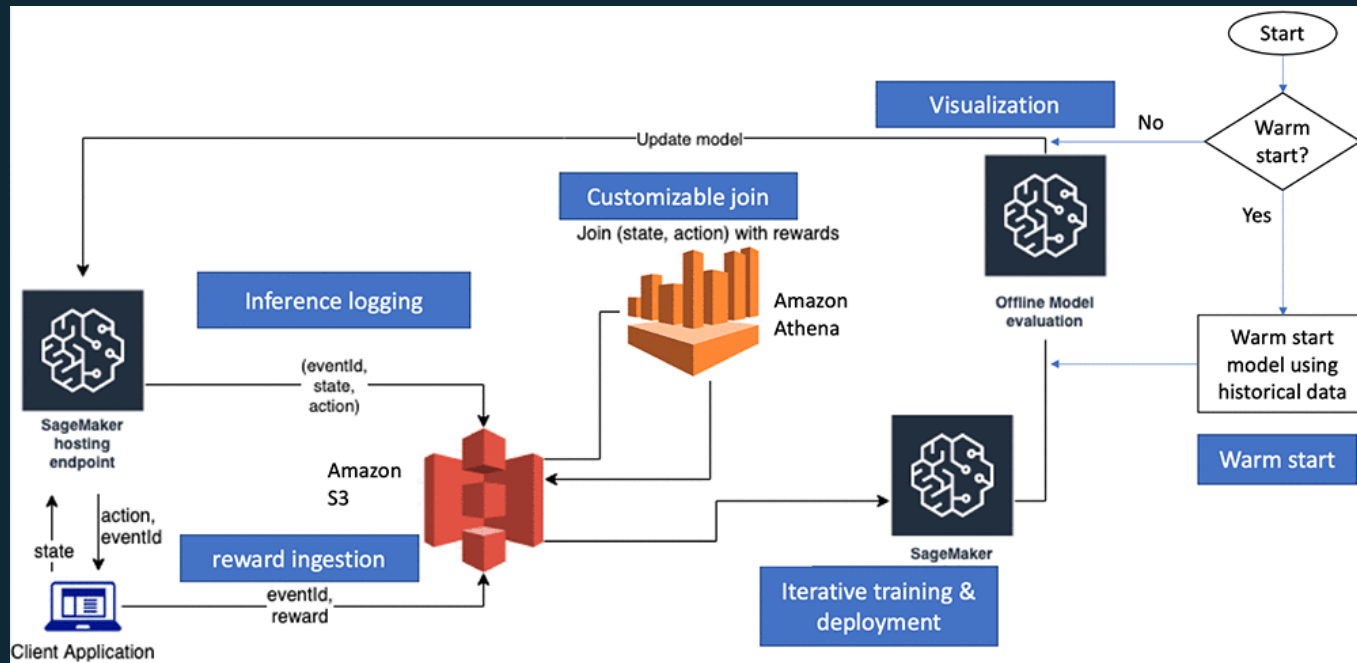
[https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/reinforcement\\_learning/rl\\_deepracer\\_robomaker\\_coach\\_gazebo](https://github.com/aws-labs/amazon-sagemaker-examples/tree/master/reinforcement_learning/rl_deepracer_robomaker_coach_gazebo)

# Let's look at some other examples

*These examples and many others are available today in Amazon SageMaker RL*


# Amazon SageMaker RL contextual bandits solution

personalized web services (content layout, ads, search, product recommendations, etc.) are continuously faced with decisions to make, often based on some contextual information






# Autoscaling



Objective	Adapt instance capacity to load profile
State	Current load, failed jobs, active machines
Action	Remove or add machines
Reward	Positive for successful transactions
	High penalty for losing transactions

<https://aws.amazon.com/blogs/aws/amazon-sagemaker-rl-managed-reinforcement-learning-with-amazon-sagemaker/>

# Financial portfolio management



Objective	Maximize the value of a financial portfolio
State	Current stock portfolio, price history
Action	Buy, Sell stocks
Reward	Positive when return is positive
	Negative when return is negative

[https://github.com/awslabs/amazon-sagemaker-examples/tree/master/reinforcement\\_learning/rl\\_portfolio\\_management\\_coach\\_customEnv](https://github.com/awslabs/amazon-sagemaker-examples/tree/master/reinforcement_learning/rl_portfolio_management_coach_customEnv)

***Thank You***